



Open in app

Get started



Published in Towards Data Science

You have **2** free member-only stories left this month. [Sign up for Medium and get an extra one](#)




Sadrach Pierre, Ph.D.

Follow

Aug 20, 2020 · 6 min read ★

# Predicting Missing Values with Python

## Building Models for Data Imputation

[Source](#)

For data scientists, handling missing data is an important part of the data cleaning and





Open in app

Get started

Let's get started!

For our purposes, we will be working with the wines dataset which can be found [here](#).

To start, let's read the data into a Pandas data frame:

```
import pandas as pd
df = pd.read_csv("winemag-data-130k-v2.csv")
```

Next, let's print the first five rows of data:

```
print(df.head())
```

```

country                description \
0      Italy  Aromas include tropical fruit, broom, brimston...
1  Portugal  This is ripe and fruity, a wine that is smooth...
2         US  Tart and snappy, the flavors of lime flesh and...
3         US  Pineapple rind, lemon pith and orange blossom ...
4         US  Much like the regular bottling from 2012, this...

              designation  points  price  province \
0          Vulkà Bianco      87    NaN  Sicily & Sardinia
1          Avidagos        87    15.0        Douro
2              NaN        87    14.0        Oregon
3  Reserve Late Harvest      87    13.0      Michigan
4  Vintner's Reserve Wild Child Block  87    65.0        Oregon

              region_1      region_2  taster_name \
0              Etna            NaN  Kerin O'Keefe
1              NaN            NaN    Roger Voss
2  Willamette Valley  Willamette Valley  Paul Gregutt
3  Lake Michigan Shore            NaN  Alexander Peartree
4  Willamette Valley  Willamette Valley  Paul Gregutt

taster_twitter_handle                                     title \
0      @kerinokeefe  Nicosia 2013 Vulkà Bianco (Etna)
1      @vossroger   Quinta dos Avidagos 2011 Avidagos Red (Douro)
2      @paulgwine   Rainstorm 2013 Pinot Gris (Willamette Valley)
```





Open in app

Get started

0	White Blend	Nicosia
1	Portuguese Red	Quinta dos Avidagos
2	Pinot Gris	Rainstorm
3	Riesling	St. Julian
4	Pinot Noir	Sweet Cheeks

Let's take a random sample of 500 records from this data. This will help with speeding up model training and testing, though it can easily be modified by the reader:

```
import pandas as pd
df = pd.read_csv("winemag-data-130k-v2.csv").sample(n=500,
random_state = 42)
```

Now, let's print the info corresponding to our data which will give us an idea of which columns have missing values:

```
print(df.info())
```





Open in app

Get started

#	Column	Non-Null Count	Dtype
0	country	500 non-null	object
1	description	500 non-null	object
2	designation	369 non-null	object
3	points	500 non-null	int64
4	price	471 non-null	float64
5	province	500 non-null	object
6	region_1	424 non-null	object
7	region_2	188 non-null	object
8	taster_name	407 non-null	object
9	taster_twitter_handle	385 non-null	object
10	title	500 non-null	object
11	variety	500 non-null	object
12	winery	500 non-null	object

Several columns have less than 500 non-null values, which correspond to missing values. First let's consider building a model that imputes missing 'price' values using the 'points'. To start, let's print the correlation between 'price' and 'points':

```
print("Correlation: ", df['points'].corr(df['price']))
```

**Correlation: 0.4161667418606222**

We see that there is a weak positive correlation. Let's build a linear regression model that uses 'points' to predict the 'price'. First, let's import the 'LinearRegression' module from 'scikit-learn':

```
from sklearn.linear_model import LinearRegression
```



[Open in app](#)[Get started](#)

filter out the missing values by selecting only positive price values:

```
import numpy as np
df_filter = df[df['price'] > 0].copy()
```

Let's also initialize lists we will use to store our predictions and actual values:

```
y_pred = []
y_true = []
```

We will use K-fold cross validation to validate our model. Let's import the 'KFolds' module from 'scikit-learn'. We will use 10 folds to validate our model:

```
from sklearn.model_selection import KFold
kf = KFold(n_splits=10, random_state = 42)
for train_index, test_index in kf.split(df_filter):
    df_test = df_filter.iloc[test_index]
    df_train = df_filter.iloc[train_index]
```

We can now define our input and output:

```
for train_index, test_index in kf.split(df_filter):
    ...
    X_train = np.array(df_train['points']).reshape(-1, 1)
    y_train = np.array(df_train['price']).reshape(-1, 1)
    X_test = np.array(df_test['points']).reshape(-1, 1)
    y_test = np.array(df_test['price']).reshape(-1, 1)
```

And fit our linear regression model:





Open in app

Get started

```
model = LinearRegression()
model.fit(X_train, y_train)
```

Now let's generate and store our predictions:

```
for train_index, test_index in kf.split(df_filter):
    ...
    y_pred.append(model.predict(X_test)[0])
    y_true.append(y_test[0])
```

Now let's evaluate the performance of our model. Let's use mean squared error to evaluate the performance of our model:

```
print("Mean Square Error: ", mean_squared_error(y_true, y_pred))
```

**Mean Square Error: 590.0704441239659**

We see that the performance isn't too great. We can improve this by training on prices bound by the mean price plus one standard deviation:

```
df_filter = df[df['price'] <= df['price'].mean() + df['price'].std()
].copy()
...
print("Mean Square Error: ", mean_squared_error(y_true, y_pred))
```

**Mean Square Error: 83.31517408259938**

While this significantly improves performance this comes at the price of not being able to accurately impute values for highly priced wines. Instead of using a regression model of a



[Open in app](#)[Get started](#)

wine 'price'. First, let's convert the categorical variables into categorical codes that can be handled by the random forests model:

```
df['country_cat'] = df['country'].astype('category')
df['country_cat'] = df['country_cat'].cat.codes

df['province_cat'] = df['province'].astype('category')
df['province_cat'] = df['province_cat'].cat.codes

df['winery_cat'] = df['winery'].astype('category')
df['winery_cat'] = df['winery_cat'].cat.codes

df['variety_cat'] = df['variety'].astype('category')
df['variety_cat'] = df['variety_cat'].cat.codes
```

Let's increase the random sample size to 5000 :

```
df = pd.read_csv("winemag-data-130k-v2.csv").sample(n=5000,
random_state = 42)
```

Next, let's import the random forest regressor module from scikit-learn. Let's also define the list of features we will use to train our model:

```
from sklearn.ensemble import RandomForestRegressor
features = ['points', 'country_cat', 'province_cat', 'winery_cat',
'variety_cat']
```

Let's train our model using a random forest with 1000 estimators and a max depth of 1000. Let's then generate predictions and append them to a new list:

```
for train_index, test_index in kf.split(df_filter):
```





Open in app

Get started

```
y_test = np.array(df_test['price'])
model = RandomForestRegressor(n_estimators = 1000, max_depth =
1000, random_state = 42)
model.fit(X_train, y_train)

y_pred_rf.append(model.predict(X_test)[0])
y_true_rf.append(y_test[0])
```

Finally, let's evaluate the mean squared error for both the random forest and the linear regression models:

```
print("Mean Square Error (Linear Regression): ",
mean_squared_error(y_true, y_pred))
print("Mean Square Error (Random Forest): ",
mean_squared_error(y_pred_rf, y_true_rf))
```

```
Mean Square Error (Linear Regression): 104.60465409290084
Mean Square Error (Random Forest): 40.199467115694446
```

We see that the random forests model has superior performance. Now, let's predict the missing price values using our models and display sample predictions:

```
df_missing = df[df['price'].isnull()].copy()

X_test_lr = np.array(df_missing['points']).reshape(-1, 1)
X_test_rf = np.array(df_missing[features])

X_train_lr = np.array(df_filter['points']).reshape(-1, 1)
y_train_lr = np.array(df_filter['price']).reshape(-1, 1)

X_train_rf = np.array(df_filter[features])
y_train_rf = np.array(df_filter['price'])

model_lr = LinearRegression()
model_lr.fit(X_train_lr, y_train_lr)
```







Open in app

Get started

```
print("Random forests regression predictions: ",  
      model_rf.predict(X_test_rf)[0])
```

**Linear regression predictions: 28.05582905629467**  
**Random forests regression predictions: 24.94**

I'll stop here but I encourage you to play around with feature selection and hyperparameter tuning to see if you can improve performance. Further, I encourage you to extend this data imputation model to impute missing values in categorical fields such as 'region\_1' and 'designation'. Here you can build a tree-based classification model trained on categorical and numerical features to predict the missing values for the categories listed.

## CONCLUSIONS

To summarize, in this post we discussed how to build machine learning models that we can use to impute missing values in data. First, we built a linear regression model trained on 'points' for reviewed wines to predict the price of wines. We then built a random forest model trained on 'points' and additional categorical variables to predict wine prices. We saw that the random forests model significantly outperformed the linear regression based data imputation model. I hope you found this post useful/interesting. The code from this post is available on [GitHub](#). Thank you for reading!

---

## Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and





Open in app

Get started

